

Master of Science in Informatics at Grenoble
Master Informatique
Specialization Graphics, Vision and Robotics

Procedural Stylization

Isnel Maxime

June 2019

Research project performed at INRIA

Under the supervision of:
Romain Vergne
Joëlle Thollot

Defended before a jury composed of:
James Crowley
Dominique Vaufreydaz

June

2019

Abstract

Stylization is the action of representing an object in 3 dimensions in a simple and decorative way. This 3D scene stylization uses images or filtering. It is often subject to temporal coherence problems. These problems are the subject of this report.

Temporal coherence problems are present in animated scenes. For example, elements of the image appear and disappear in the next image, brushstrokes do not follow the movement of the object, pencil strokes become larger as you get closer to the object, etc.

In this report, we describe a method that aims to minimize these temporal consistency problems while maintaining control over the stylization of the object. To do this, our method uses the principles of stylization in an image based on the collage of images containing brushstrokes, hair, dots, leaves, etc. We also use 3D textures to control the movement of these images on the screen.

Résumé

Styliser est l'action de représenter un objet en 3 dimensions de manière simple et décorative. Cette stylisation de scène 3D utilise des images ou du filtrage. Elle est souvent sujette à des problèmes de cohérence temporelle. Ces problèmes font l'objet de ce rapport.

Les problèmes de cohérence temporelle sont présents dans les scènes animées. Par exemple, des éléments de l'image apparaissent et disparaissent à l'image suivante, les coups de pinceau ne suivent pas le mouvement de l'objet, les coups de crayon deviennent plus gros quand on se rapproche de l'objet, etc.

Dans ce rapport, nous décrivons une méthode qui vise à minimiser ces problèmes de cohérence temporelle tout en gardant le contrôle sur la stylisation de l'objet. Pour cela, notre méthode utilise les principes de stylisation dans une image à base de collage d'images contenant des coups de pinceau, des cheveux, des points, des feuilles, etc. Nous utilisons également des textures 3D afin de contrôler le mouvement de ces images sur l'écran.

Acknowledgement

Avant tout développement, il me semble important de remercier toutes les personnes ayant contribué au bon déroulement de mon stage et qui ont rendu cette expérience riche, variée et unique.

Je voudrais tout d'abord remercier mes tuteurs de stages, Joëlle Thollot et Romain Vergne, pour leur patience, leurs explications et leur disponibilité qu'ils m'ont accordé. Ils m'ont partagé leurs savoirs sur le rendu non-photoréaliste et rendu graphique, leurs compétences de gestion d'un projet de recherche et leurs connaissances techniques.

Je voudrais aussi remercier toute l'équipe MAVERICK pour leur chaleureux accueil. Il fut un réel plaisir de travailler avec toutes ces personnes attentionnées et il fut plaisant et flatteur d'être considérée comme un membre de cette équipe. Pour finir, je voudrais remercier l'INRIA pour m'avoir procuré les ressources nécessaires pour mener à bien ce stage et cette recherche.

Contents

Abstract	i
Résumé	i
Acknowledgement	i
1 Introduction	1
1.1 Background	1
1.2 Problem Statement	1
1.2.1 Flatness	2
1.2.2 Motion Coherence	2
1.2.3 Temporal continuity	2
1.3 Contents of this report	3
2 Previous Work	5
2.1 Object Space	5
2.2 Image space	8
3 Realisation	13
3.1 Overview	13
3.2 Procedural noise and fractalization	14
3.3 Splatting	18
3.4 Stylization	19
4 Practical implementation	23
5 Results and performance	27
6 Discussion and conclusion	31
6.0.1 Discussion	31
6.0.2 Conclusion	31
Bibliography	33

Introduction

1.1 Background

A part of computer graphics is to create non-photorealistic images. A method to do this is to stylize 3D scenes. Stylizing an object means creating an image that imitates the style of an artist who would have drawn it on a sheet of paper. There exist many different styles like hand drawing, brush painting, pointillism painting, stippling, watercolor painting, etc. The two main methods to stylize is to use textures mapped on the objects of the scene or to use marks/splats which are 2d images.

1.2 Problem Statement

The main problem of stylizing a 3D object in an animation is *temporal coherence*. The *temporal coherence* problem in non-photorealistic rendering encompasses both spatial and temporal aspects of the marks. The effect given by the stylization has to be kept if the object is moving, rotating and scaling. Many research has been done to solve this problem of *temporal coherence* [29, 7, 3]. Bénard et al. separate this problem in three sections inspired by previous work[19, 11, 5, 8] but they show that an ideal solution (Figure 1.1 a) does not exists. Even *Loving Vincent* [2] that was done by artists frame by frame has problems of marks popping because they cannot follow perfectly the flow motion of a scene. Neglect one of these three goals provide artifacts (Figure 1.1 b-d).

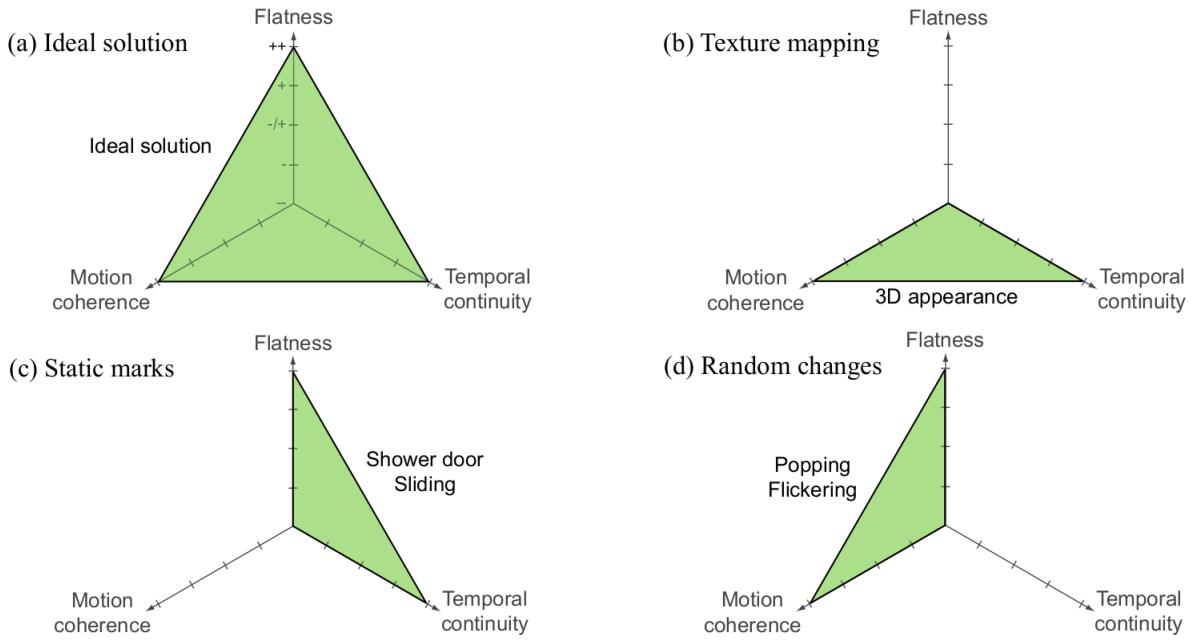


Figure 1.1: Problem involve by temporal coherence depending of the flatness, temporal continuity, motion coherence.

1.2.1 Flatness

The impression of drawing on a flat surface gives the *flatness*. The stylization has a good *flatness* if the image rendered has a good 2D appearance. In order to keep this effect, the size and the distribution of the marks of your stylization have to be independent of the distance between the stylized object and the camera.

1.2.2 Motion Coherence

Motion coherence is a correlation between the motion of marks and the motion of the 3D object. Bad *Motion coherence* will give the impression to see the scene through a semi-transparent layer of marks, this is called *shower door* effect [19], an example to illustrate what happens when there is a bad *Motion coherence* is the movie *Loving Vincent*[2]. The goal is to provide in 2D screen space a perceptual impression of motion as close as possible to the 3D displacement in object space.

1.2.3 Temporal continuity

Temporal continuity is the quality of minimizing changes from frame to frame to ensure fluid animations. In order to have good *temporal continuity*, the marks of the image have to fade slowly during the animation. Human perception is very sensitive to *temporal incoherence* according to some perceptual studies[30, 26].

The problem introduced in the works of Bénard et al.[8] is the difficulty to have a ideal solution, the one that has a good *flatness*, a good *motion coherence* and a good *temporal continuity*. These three goals are inherently contradictory when you improve one you neglect one or maybe more. So researchers work to find solutions that make *trade-offs* between these three goals. We built our approach trying to take the better of each different approach, looking at what works well and what can be integrated into our method.

1.3 Contents of this report

After an analysis of the previous works on stylizing scene, we searched on how we could have a solution that permits to control the 3 axes of the *temporal coherence*. We defined what were the advantages of each method of stylizing: marks-based method and texture-based method. The texture-based method has a very good *temporal continuity* and *motion coherence* at a price of bad *flatness* and a poor variety of style. On the over hand, the mark-based method has good flatness and a wide variety of possibles styles but introducing a problem of anchoring these marks. From this analysis, we decided to use marks to stylize our object but we had to find a way to anchor these splat. We used the good property of the textures mapped in object space. They permit to have a perfect *temporal continuity* and *motion coherence* so we used the textures to anchor our marks to the object but we did not ordinary texture to do so, we used procedural textures because we can manipulate them easily and are cheap to render.

This is our solution, draw marks anchored to the object thanks to procedural textures in an image. We implemented this solution in OpenGL in order to render all images with the power of the GPU. In the figure 1.2 you can see some results done with our method.



Figure 1.2: Some results of stylizing with our method of procedural stylization (left: pointillism rendering, middle: brush painting, right: hair rendering).

— 2 —

Previous Work

Many research have been done in stylizing 3D scenes[25, 24, 16, 7, 6, 13, 8] trying to propose solutions or trade-offs for the problem of *temporal coherence*. In this section of this report, we will present you techniques to stylize 3D scenes and we will show their advantages and disadvantages. In order to render an image of a 3D object in the screen, a graphics program goes through several steps that compute some different information like the gradient of the image, the shadows made by the object, the amount of light received by the object, etc. The gathering of all these steps is called *graphical pipeline rendering*. In this graphical pipeline, there are two locations where we can stylize the objects. The process of stylizing can be applied either on the surface of the 3D objects (*object-space*) or on the pixels of the screen after the models have been rendered (*image-space*). We will treat these two space separately and with the two different types of methods to stylize.

2.1 Object Space



Figure 2.1: Texture mapping: example

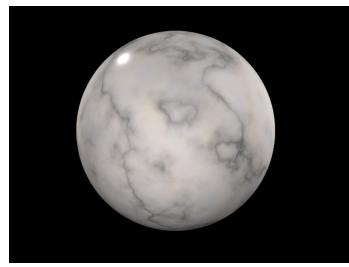


Figure 2.2: Marble with procedural texture

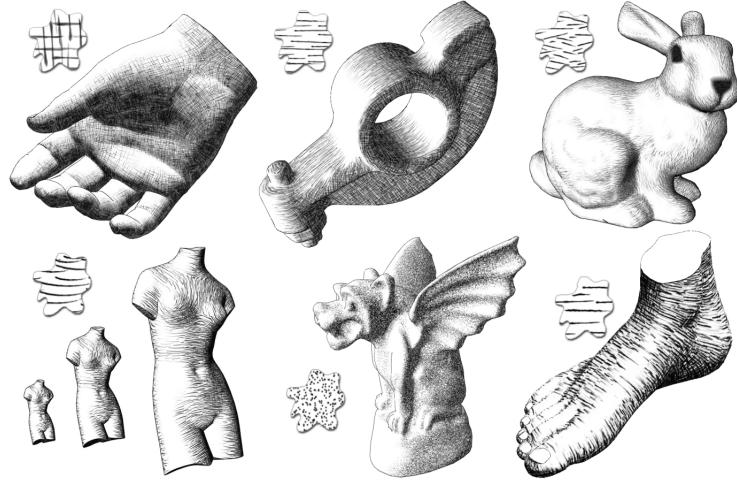


Figure 2.3: Real time hatching rendering [24]



Figure 2.4: Watercolor stylization with procedural textures [7]

In the object space, we work on the surface of the object and so we have all the knowledge about the geometry.

Texture-based methods

One of the most used ways to colored object in 3D is the *texture mapping*. It consists of mapping an image on the object (Figure 2.12.3). This technique is widely used in video games because it is easy to implement, it can be implemented for GPU and it needs low computation. As shown in the example, texture mapping can be used with images as a texture in order to stylize the object[24, 16, 13]. Texture mapping can also be done with *procedural noises*[23] in this case called *procedural textures*. Procedural textures are mathematically computed from 3D surfaces coordinates the most famous is Perlin but there exist others like Gabor noise, Worley noise, etc. With these textures, we can create images looks like marble (Figure 2.2) or realist wood. Bénard et al.[7, 6] use this method of texture mapping with noise in order to make watercolor stylization (Figure 2.4).

This method has a problem of *foreshortening* on the silhouettes which makes an area with more elements than in the center of the object and when zooming in/out the size of the elements in the mapped texture vary depending on the distance from the camera, if we get closer to the object the elements on the texture become bigger and bigger going up to pixelization sometimes. These two problems make texture mapping bad in term of *flatness* because an artist does not draw bigger strokes if the object is close and does not draw more details at the silhouettes than in the center of his objects. He is constraint by the size and shape of the tools that we use. Texture mapping is also bad in term of variety of style that you can do because you are depending on the texture and chosen you cannot, for instance, do a rendered image that looks like painted with brushes and especially you cannot change the shape of the object in the rendered image, a perfect sphere will always appear as a perfect sphere. On the other hand, this method naturally ensures *motion coherence* and *temporal continuity* because the texture is mapped directly on the surface of the object.



Figure 2.5: OverCoat: implicit canvas [25]

Mark based methods

The natural way to stylize 3D objects is to as an artist apply paint strokes on the object. These paint strokes can be represented with smalls images also called splat. Some works [19, 1, 10](more in the state of the art [8]) use point distribution in order to make anchor points for splats. These point distributions are often computed in image space and then are projected on the model. Anchor these splats to the model improve the *motion coherence* because each splat will follow the motion of the 3D model. These splats are rendered in the image space as a 2D sprites so preserved the *flatness*. The problem is how to have the point distribution and

how can we control it in order to have a uniform, not too sparse and not too dense distribution. Moreover, these point distribution does not provide control over the *temporal continuity*. In our method, we use procedural noise to anchor the splats.

Daniels[12] and Schmid[25] propose to project splat on an inflated proxy geometry and stored them in nodes but anchored to the geometry of the model but this technique is expensive in term of storage. OverCoat[25] is an interactive software that permit to the user to stylize manually a 3D object (Overcoat demo). Their techniques are the closer techniques to our proposed solution. In their solution, they use 3D splines drawn by the user and they display images along these lines. Each spline is stored in nodes that contains the image to splat along this line, the color, the thickness and the transparency that permit the user to edit the stylization after drawing the splines. These splines are anchored on a proxy generated by the program this has the disadvantage of placing the stylization at the beginning of the graphical pipeline. It has a problem of *foreshortening* of the silhouettes that make the rendering blurred at some place of the image. Overcoat has a big problem of flatness when rescaling the stylized object, this is because the user draws marks at a specific scale so when we zoom in the scene the marks looks bigger or smaller and that affect the *flatness* of the rendered image. All the marks drawn are anchored to the geometry so it ensures the *motion coherence* and because they use a proxy geometry the shape of the object can be changed, some details can be added. If the object is animated it will create a hole or gathering of marks.

The mark based methods have a problem of foreshortening has said before and the non-constant size of the marks when rescaling of the object make the impression of bad *flatness*. But the *motion coherence* is ensured by anchoring the marks to geometry that makes them follow the motion of the stylized object.

2.2 Image space

In image space, we work on images with 2d coordinates that means working on the pixel of the image directly. In a normal pipeline rendering, the image space correspond to the *compositing* which is the moment when many information have already be rendered and are stored in images (like shadow map, image filter, global illumination, etc.), during the *compositing* we combine the computed images to render the final image.

Texture-based methods



Figure 2.6: Image filter for watercolorization [20]

Many methods to stylize in image space used texture based approaches. It consists to apply the texture to the entire image (Figure 2.6) [8, 20] but in the case of stylizing animated scenes, the problem is how do we deform the texture to minimize the apparition of sliding artefacts. Sliding artefacts are some parts of the image that do not move according to scene motion and give the impression that they are sliding. We can distinguish two families of approaches to solve this problem. The first family of approaches use an approximation of the 3D camera motion with 2D transformations of the texture[11]. This gives a nice trade-off between *motion coherence* and *flatness* but it is limited to static scenes and a set of few camera motions. Moreover, sliding artefacts still occur with strong parallax so Fung et al.[14] and Breslav et al.[5] improve the approximation of the scene motion in order to reduce sliding artefacts.

The second family of approaches use deformations to animate the texture[4]. These deformations are computed from the optical flow of a video. This is an extension of the methods used in vector field visualization by Neyret[22]. These deformations can distort the texture and alter the original pattern. The method of Bousseau et al.[4] is very effective with stochastic textures as the fractalization process but creates artefacts with structured patterns.

These texture based methods for 3d scenes have issues with *motion coherence* and *temporal continuity* because of the independence of each frame and also these techniques suffer to make a poor variety of style, you cannot do for example hairs, brushes painting, pencil drawings. But they can do some different styles with good *flatness*.

Mark based methods



Figure 2.7: Hertzman's algorithm on an image [27]



Figure 2.8: Example segmentation image[8]

A method very used to stylize in image space consists to draw strokes/splats at some place of the image[3, 29, 9, 31, 15]. It is like adding other images in an image (like you can do in Photoshop or Gimp, Figure 2.7). The question of these mark based method is where do we place these marks in order to have a stylized rendering without losing the meaning of the scene and without losing details. A first approach is to extract lines that are relevant like the silhouettes, etc. [29, 15, 17] and then stylize the image with this information, like keeping only the extracted lines and change the shape of each line or apply strokes along these lines as Vergne et al.[29] did try to have a good *temporal coherence*. The problem of these techniques is the popping marks due to a bad *temporal continuity* because we do not know if in the next frame there will be more details or not or the color to display is different from the current one. A second approach is to segment the image (Figure 2.8) in order to have the different parts of the scene[31, 18]. Thanks to this segmentation, they apply different strokes for each part of the image with the corresponding colors. The work of Lin et al.[18] is about videos so they use the optical flow of the videos in order to have a good *temporal coherence*. These mark based methods have a good impression of *flatness* thanks to the splatting in image space, this is something that we will use in our approach.

These methods of splatting have almost the same problem that the texture based method they suffer of bad *motion coherence* and *temporal continuity*. Temporal continuity that can be improved using a motion flow which cannot be possible in our case. In the other hand, changing the type of splat, the size and their orientation we can make many different styles, like brushes painting, pencil drawings, stippling, hairs. These methods draw 2d strokes/brushes in the image so it looks like drawn by a person on a flat surface which makes a good *flatness* impression.

The figure 2.9 is a summary of the previous works. The interesting point is mark based method can make a wider variety of styles than the texture based method. We can see that for *motion coherence* and *temporal continuity* doing the stylization in object space is better than in image space but in image space, we can do stylization that has a good *flatness* impression contrary to the object space.

	Motion coherence	Flatness	Temporal continuity	Style variation
Object space				
Texture-based methods	++	--	++	--
Mark based methods	++	+/-	+/-	+/-
Image space				
Texture-based methods	-	++	-	--
Mark based methods	-	++	--	+

Figure 2.9: Summary of trade-offs made in different approaches

— 3 —

Realisation

3.1 Overview

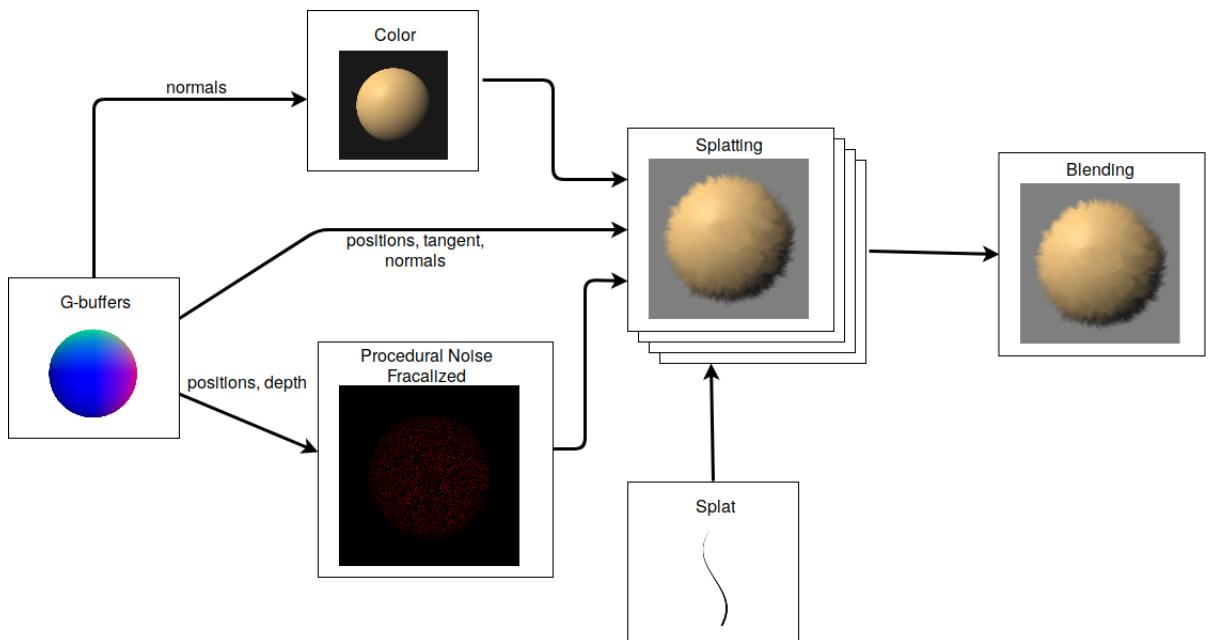


Figure 3.1: Overview of our solution.

As explained in the state of the art and in the figure 2.9 each approach has its advantages and its disadvantages. Our main idea is to use procedural texture rendered in the object space to guide the splats in image space. This solution permits to apply something like 2D images on the screen so have a good *flatness* while keeping the information on the silhouettes, the orientation, the depth, etc. This solution only needs images as input so it can easily be integrated into most of rendering pipelines without re-render images. In the figure 3.1, there is an overview of our pipeline. First we computed all the data from the G-Buffer (normals, tangents, position, matrices, depth, UV coordinates) then we render a color (in this example a simple shading) after this we computed our procedural noise that will help to anchor the splats to the object and we choose the splat that will convolve in order to make the stylization. Finally, we draw splat

on the entire screen, then we decide which one has to be displayed this is the splatting step and the last step is the blending which computed which splat is in front of the others and which color has to be prompted.

We chose to use mark based methods to stylize our scene because texture based methods in image space give a poor variety of styles as said in the work of Bénard et al.[7]. This mark based method implies to decide where in the image the splats will be drawn. In our problem, the goal is to anchor these splats with the objects in order to have the same motion for the splats and the object. This avoid the problem of *shower door effect* and ensure the good *motion coherence*. So we needed anchor points depending on the position of our object. Therefore in our approach, we used procedural noise[23] as a texture of our 3D object. The procedural noises are easy to implement, fast to compute and easy to manipulate. Like every texture computed in object space, it has a good motion coherence.

3.2 Procedural noise and fractalization

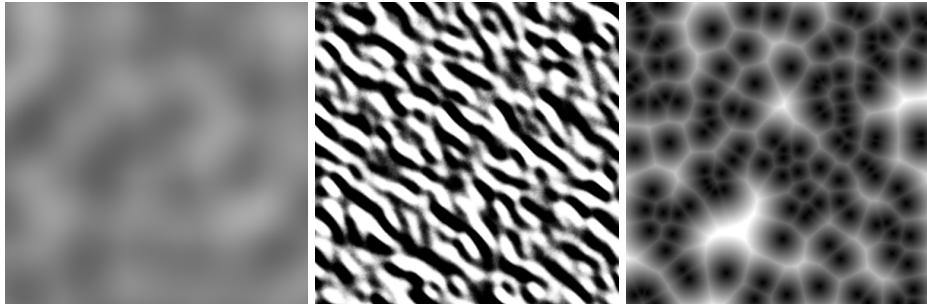


Figure 3.2: Examples of procedural texture: *left using perlin noise, middle using Gabor noise, right using worley noise.*

We compute procedural texture in order to create anchor points. In computer graphics, they are usually used to create more complex rendering like the marble texture (figure 2.2). These textures are computed from procedural noise with a mathematical process with only the position as a parameter. The algorithm to create these textures are implicit that means that with only one position you can have the value without knowing the rest of the texture. That is why they are very used in computer graphics because they can be computed in parallel on GPU. There exist much procedural noises such as Perlin noise, Worley noise, Gabor noise, Value noise, Gradient noise, etc. In order to map the procedural texture to the 3d object, we compute it with the vector position of each vertex of the object.

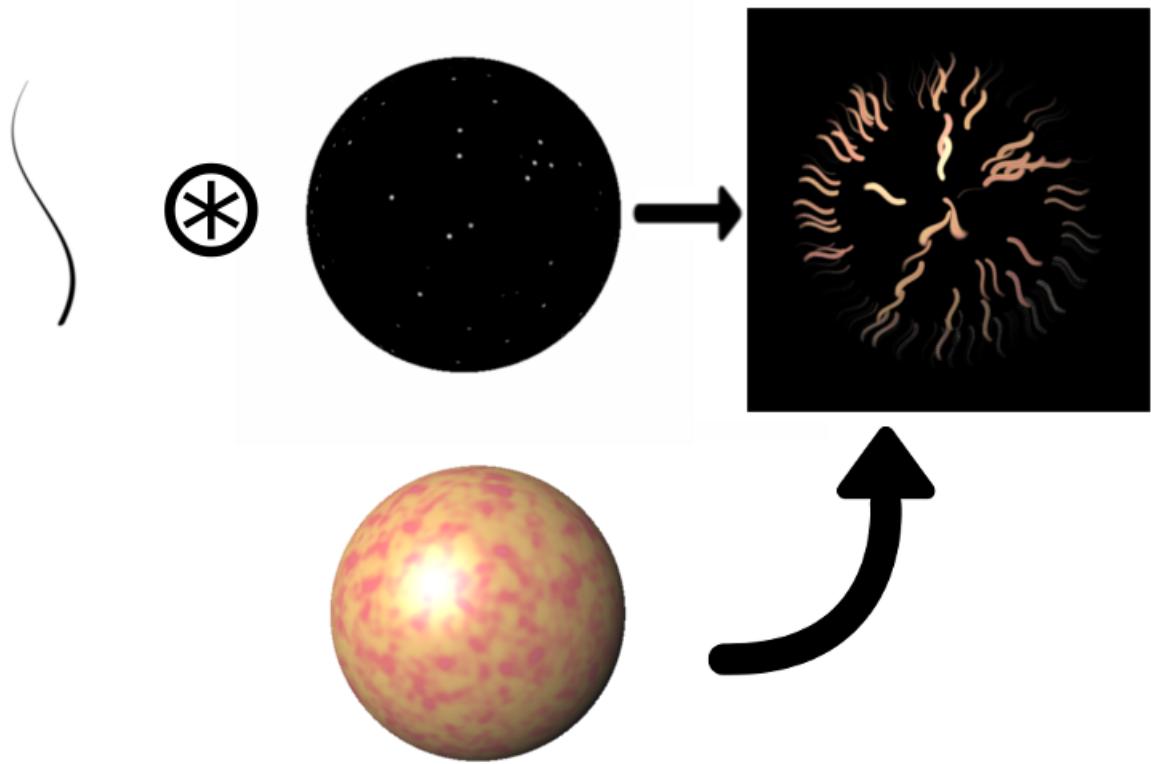


Figure 3.3: Usage of procedural texture to anchor splats: (left: splat image, middle: procedural texture, right: rendered image, bottom: color).

In our case, we used the procedural texture to anchor the splat. For each pixel of the final image we "paste" a splat. Each splat as a weight that will correspond to its opacity during the splatting and the blending. The value of this weight corresponds to the value of the noise on the current point. In the figure 3.3, you can see in the areas where the noise is close to zero there is no splat displayed in the final image. Thanks to this mechanism we can control the density of splat in the image. In our case, we worked mainly with the Worley that when inverted almost correspond to point distribution.

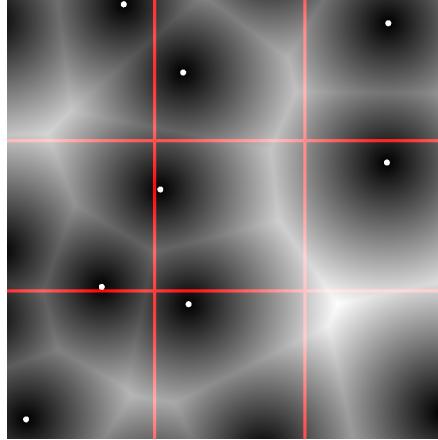


Figure 3.4: Worley noise in 2d.

We mainly work with the Worley noise which is a cellular noise as you can see at the right in the figure 3.2. To create a procedural texture from a Worley noise, we divide the image in cells (squares with the same size in 2d, red lines in the figure 3.4), then in each square we compute a point with a seeded random (white dot in the figure 3.4) and finally at each pixel of the texture we put as value the distance with the closest point computed previously with the random seed. Increase the frequency of this noise correspond to divide the image with smaller cells and so have more cells in the image. We explained the principle for a 2d texture but it is the same principle in 3d with a volumic grid. In our pipeline, we use the inverse of this noise (1-worleynoise) in order to have the white pixels of the procedural texture as anchor points. We add a threshold parameter to compute the noise in order to reduce the number of splat to display but especially to have small points in the procedural texture if not the splats convolve as you can see in the rendered image in the figure 3.3.

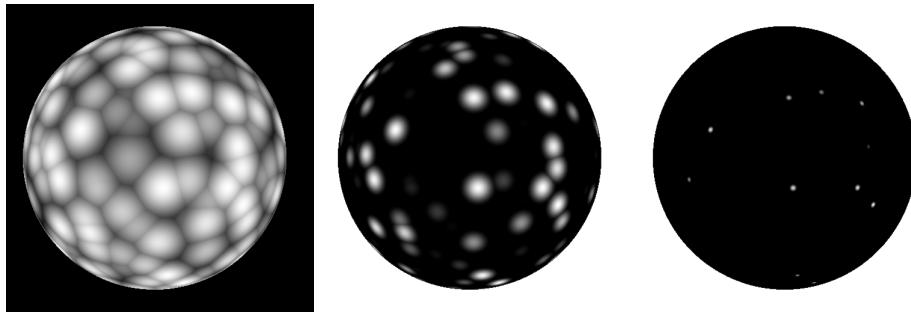


Figure 3.5: Use of threshold in the case of Worley noise with the same frequency.

As we said before we use a threshold for the purpose of having another control on the noise. With the Worley noise, this threshold corresponds to the "size" of the points. If we increase the threshold the points become smaller as you can see in the figure 3.5. In order to facilitate the control of this noise, we adapt this threshold according to the frequency and the distance from the camera. That means that the size of each point is constant if the frequency or/and the

distance from the camera is changed.

Fractalization

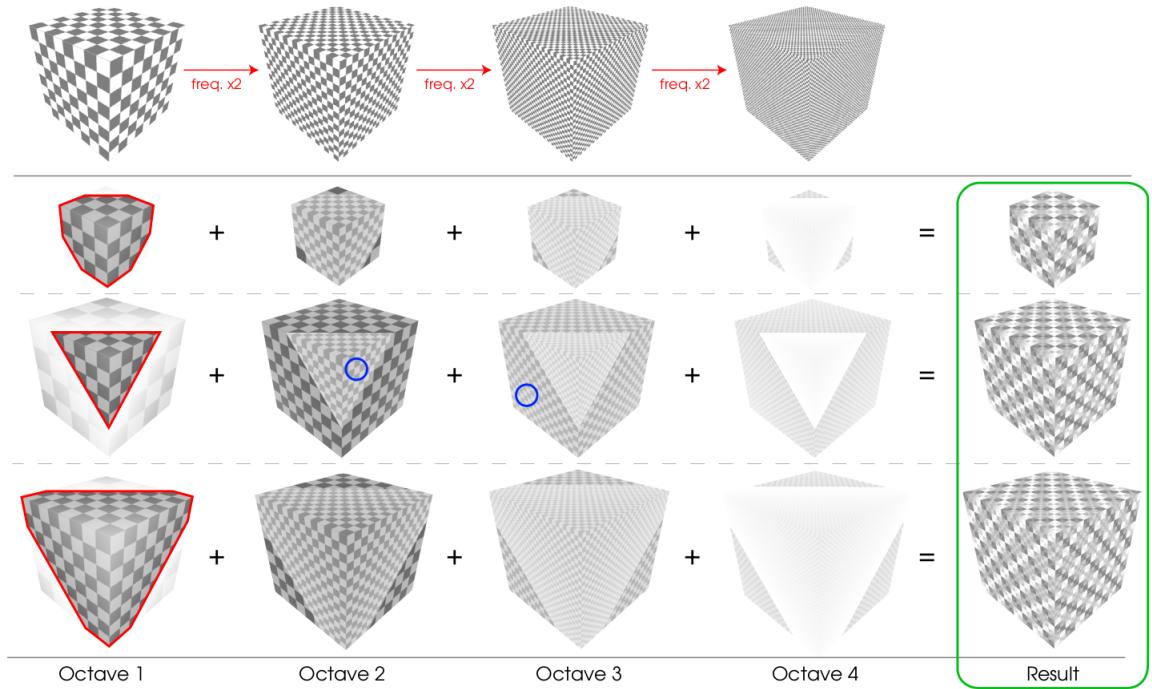


Figure 3.6: Principle of fractalization [6].

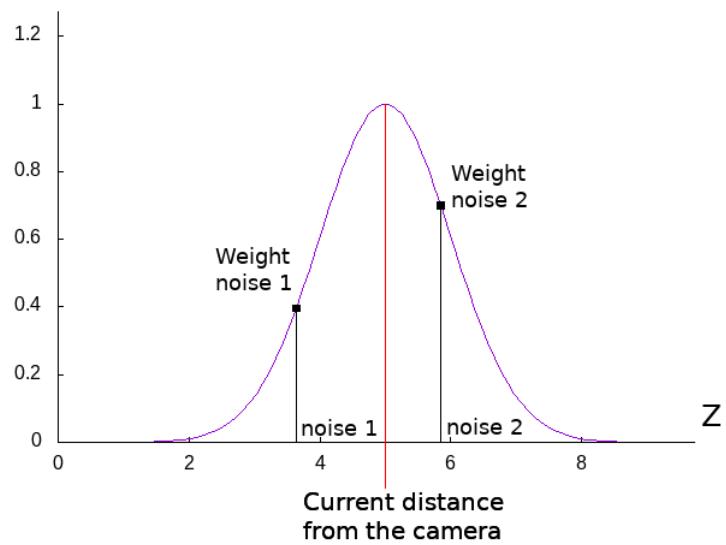


Figure 3.7: Combination of the noises with the weight value in this case with 2 noises.

Bénard et al.[6] introduce the mechanism of fractalization on textures. This technique creates an impression of infinite zoom effect (like in this example: ShaderToy of Neyret Fabrice). This is helpful in our case because when the camera is getting closer and closer or further and further less or more points in the texture. With this mechanism, we can get a quasi-constant number of points on the screen. In our method of stylizing it improves the *temporal continuity* because there is always display even if you get very close to the object and it also improves the *flatness* impression because there is almost the same number of anchor points in the screen and their size is quasi-constant. This method alters the original pattern of the texture as you can see in the figure 3.6 it can be an issue because some pattern cannot be fractalized (like the checkboard pattern) but it works well with stochastic textures. To creates this fractalization, they use textures at multiple frequencies (not necessary procedural texture) (see figure 3.6) and they combine them with the transparency and they overlap them according to the distance from the camera as you can see in the results of the figure 3.6.

The figure 3.6 show 4 noises (octaves) at different frequency it is a zooming cycle. As you can see, if the distance from the camera is close to 0 the frequency of the noises is high and if we move away from the object the frequency decrease. So when the distance from the camera is twice as large the frequency of the noise displayed is multiplied by 2. During the zoom, we modulate each octave with a weight. To compute the weight of each noise we use Gaussian interpolation centered on our current distance from the camera (see figure 3.7) with a sigma depending on the number of noises that we want to combine.

3.3 Splatting

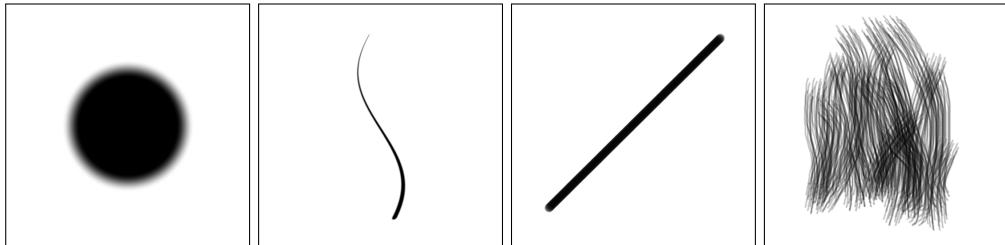


Figure 3.8: Example of what the splats can be.

The marks based method try to imitate how the artists do when they stylize. They draw on a flat surface that gives a good impression of *flatness*. We use the same principle to stylize our scene. We put splats/marks directly on the image like an artist will do (illustrated in the figure 3.3).

In our solution, the user controls the splat image he can use anything. These marks can be leaves, hairs, dots, feathers, lines, paint brushes, etc. They can also be created from a procedural noise. The user has also the control on the size of these splats, he has global control of the size and control on a specific part of the object through a texture. The rotation of the splats is also a parameter that the user can control. He can choose to rotate all the splat, for example,

depending on the normals or depending on the tangents.

Artists control how their marks are combined for example if he is doing painting he could want to have non-transparent marks in order to cover the marks behind the new one. He could want more transparency if for example, he wants to do watercolorization.

In our method, we take this into account during the blending of all marks. The user can control how many marks the stylization will mix and control the transparency (alpha blending). For example, the user can choose to only show the marks at the top of the flat surface (the closest to the camera) or he can choose to mix some marks in order to account some marks covered.

3.4 Stylization

As explained in the previous section, in our method, the user controls the content of the splats, their size, their orientation, the density of splat on the object and how they are blended with the alpha blending. In this section, we will show you some examples of how we did specific style using our pipeline.

Pointillism



Figure 3.9: Example of pointillism with our method.

In order to make pointillism style we use a dot as splat (gaussian function in 2d, the first in the figure 3.8), the procedural noise to control the number of splats has a high frequency with a high threshold in order to have small anchor points, the size of the splats can vary depending on which size of brush he wants to imitate but usually we use small splat size for this style and for the alpha blending it usually adapt to not mix many marks. In this case, the orientation does

not matter. The figure 3.9 shows the final rendered image.

Brush painting



Figure 3.10: Example of painting with brushes with our method.

In order to make an image made like a painter we use an image of a paint strokes (like the fourth in the figure 3.8), we do not need too much splat displayed so the frequency of the procedural noise does not have to be very high but a threshold high/medium, the size of the splats depend of the 3d model but usually we use large splat. For the alpha blending it depends on the effect wanted, in our case, we mixed many splats. And for the orientation usually the artists align the strokes with the tangents but it works well if the strokes are horizontal. The figure 3.10 shows the final rendered image.

Hair rendering



Figure 3.11: Example of hairs rendering with our method.

In order to show that we have the control of the *flatness* impression we tried to create hairs with a 3d aspect like it can be done in digital painting (see figure 3.11). To do this we use a splat created mathematically that looks like a hair (the second in the figure 3.8) rotated according to the normals of the object. With the procedural noise that is used to anchor marks, the density can be varied impacting directly the density of hairs in the fur. The size of the splats varies depending on what size you want your hair to be. And for the alpha blending, you do not want to mix too many splats but you do not want to only show the last hairs drawn. The figure 3.11 shows the final rendered image.

Here there are 3 examples of how we use the method to do some styles but the same styles can be done differently and others styles can be done as well. These 3 examples show what is possible to do with this method and show the control of the *flatness* we have.

— 4 —

Practical implementation

During our research on stylization, we used Gratin [28] which a pipeline rendering software using *OpenGL*. It permits to easily creates rendered images with *GLSL* combining previous images as textures. As we said before every texture used from our contribution can be computed in parallel on GPU including the noises, the fractalization of the noises, the alpha blending and the splatting. Our method use only images as input so it can easily be integrated into every pipeline rendering. For example, if you have a pipeline rendering with *path tracing* to compute the color of your scene you can use this rendered image as color input of our method without re-compute it.

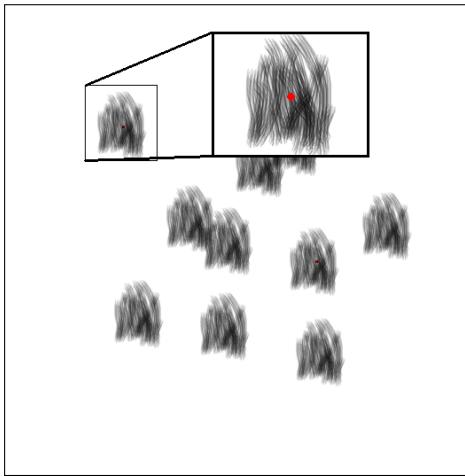


Figure 4.1: Splatting principle: red dots are the anchor points and the splat is centered and anchored to this point.

For the splatting step, we draw on each pixel of the screen a square centered on the current pixel (see figure 4.1). Every square is treated independently on the GPU before in the *vertex shader* which manipulate the 4 vertices of the square which permits the resizing of the splat and permits the rotation of the splat. After this step of resizing and rotation, we pass in the *fragment shader* which manipulate all pixels of splat. It is in this step that we will decide if we display the splat or not thanks to the procedural texture previously computed.

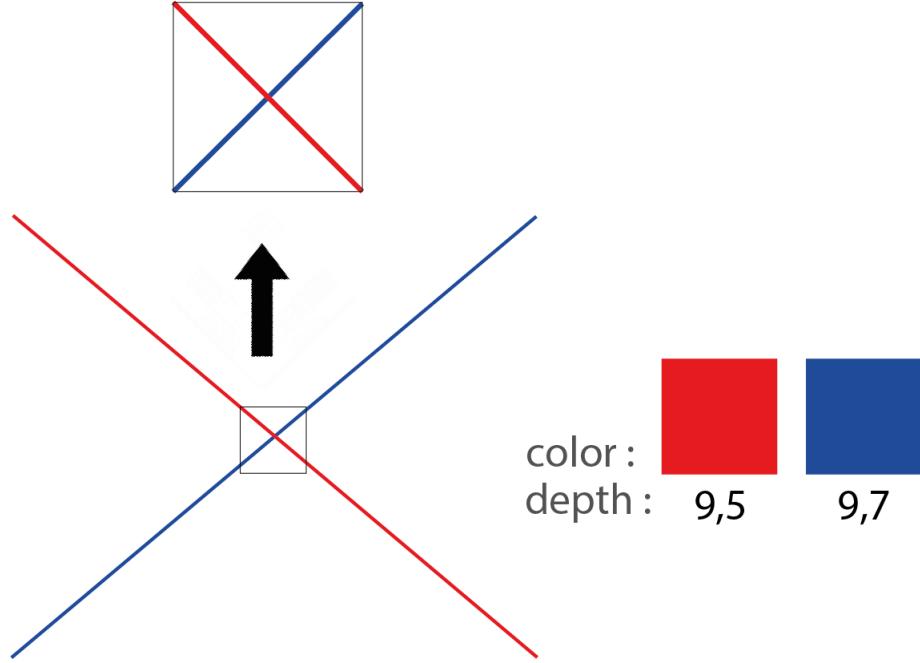


Figure 4.2: Example of usage of the order independent transparency.

In this *fragment shader*, we prepare the technique of *order-independent transparency* [21] that consists of a structure that permit to know the order of the splats depending on the depth. We create in real time for each pixel of the screen a linked list on the GPU in order to sort the pixels contained in the splat. Thanks to this, we can apply a function of alpha blending that takes account the position of the splat, their color, their opacity. In our example in the figure 4.1, we have some splats that cover others splats. How do we know which one is in front of the other if we treat them independently? For example, if there is a pixel where splat covers another one, we construct an array with 2 elements composed of color and depth so we can know which one is in front of the other one (see example in the figure 4.2). We apply the same technique but in our case we have a lot of splats that cover others and as depth to put in this array we use the depth corresponding to vertex of the stylized object where the splat is anchored (the red dots in the figure 4.1 correspond to a vertex of the 3d objects).

After creating these array, we will in the last pass treat each pixel of the screen independently and (still on GPU) decide which color has to be displayed. In order to do this, we sort the array corresponding to the current pixel according to the depth. And finally, we apply our algorithm of alpha blending (the same as the mode "over" in Overcoat [25]):

```

vec4 over(in vec4 c1,in vec4 c2) {
    float a = (c1.a+(1.-c1.a)*c2.a);
    return vec4((c1.rgb*c1.a + c2.rgb*(1.-c1.a)*c2.a)/a,a);
}

vec4 alpha_blending(vec4 colorArray[], int count, vec4
backgroundColor, float gammaBlend){
    // colorArray is the list of color sorted: colorArray[0]
    // correspond to the farest
    // count corresponds to the number of splat that we will take
    // account during the blending
    // backgroundColor is the color of the background our image
    // gammaBlend is a user parameter that will increase the
    // transparency of the displayed color

    vec4 finalColor = vec4(0,0,0,0);

    for(int i=count-1; i>=0; i--) {
        vec4 currentColor = colorArray[i];
        finalColor = over(finalColor,vec4(currentColor.rgb,clamp(pow(
            currentColor.a,1./gammaBlend),0.,1.)));
    }

    return over(finalColor,backgroundColor);
}

```

After this algorithm, we have the color to display in the current pixel of the image/screen. With all these pixels, we have our final image.

— 5 —

Results and performance

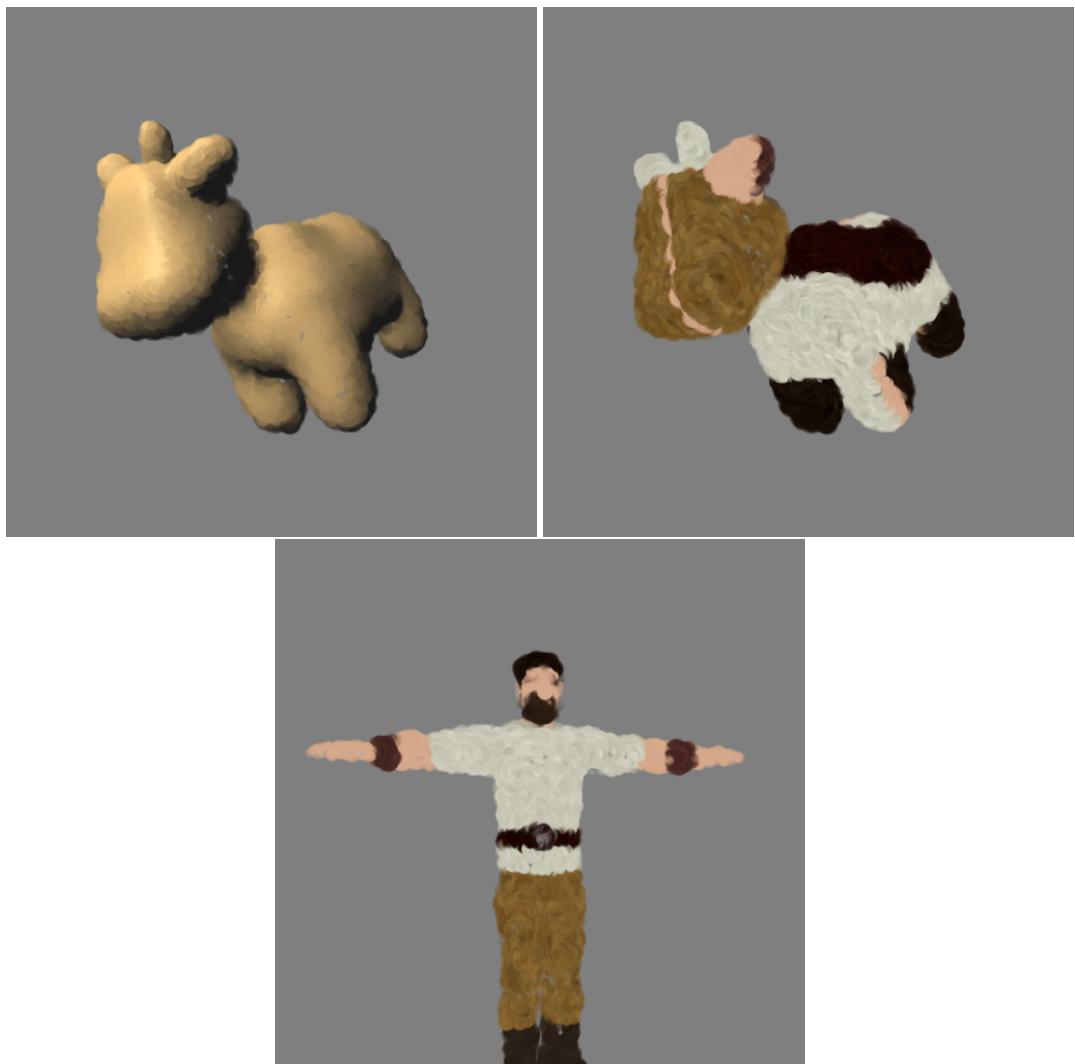


Figure 5.1: Results of pointillism.



Figure 5.2: Results of brush painting.

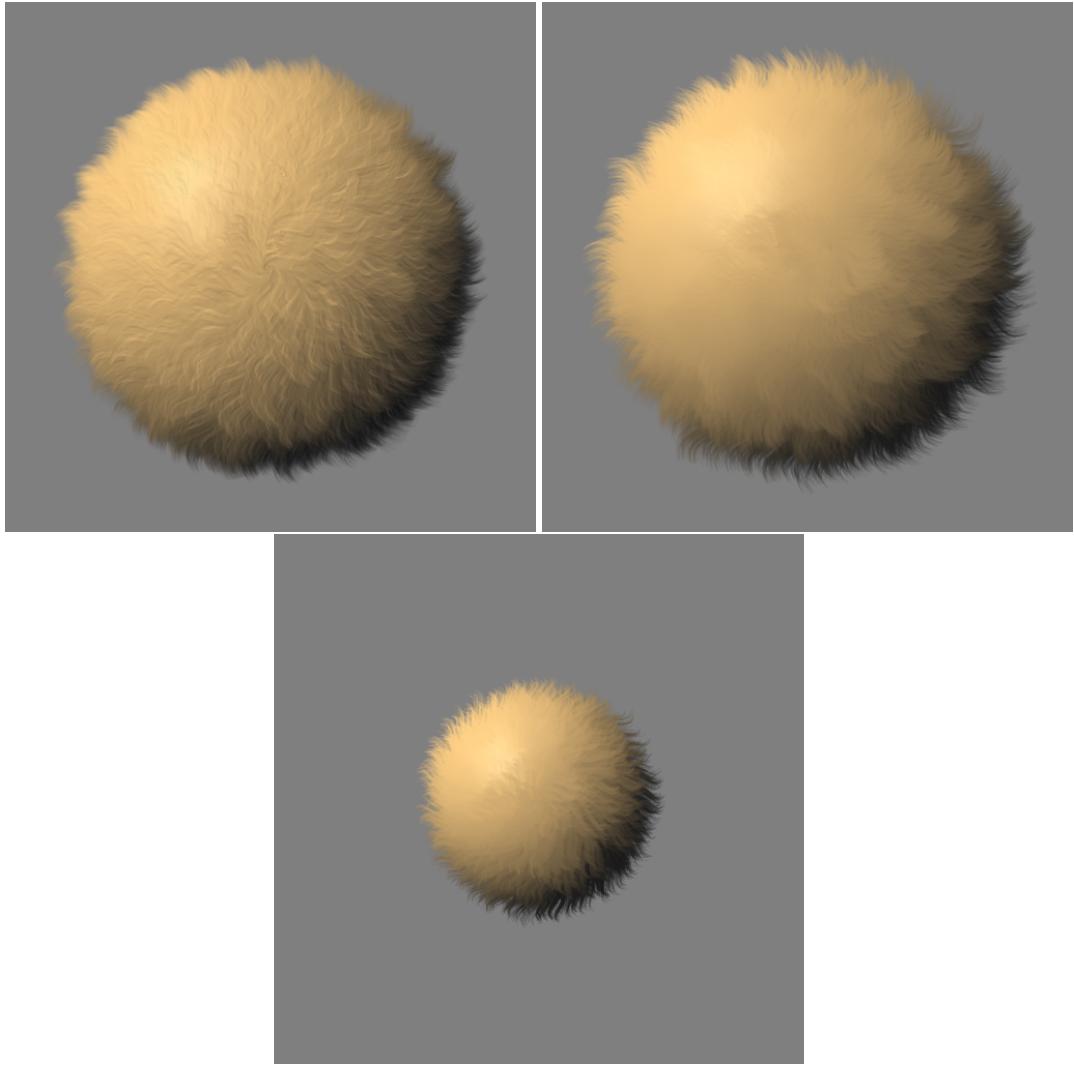


Figure 5.3: Results of hairs rendering.

Here there are some results of our method of stylization. The figure 5.1 show some rendering with a point as mark and with different colors and models. We use the dot splat (the first in the figure 3.8) with a small size in order to keep a good level of details. In the figure 5.2 we tried to stylize as painted with brushes so we used brush stroke as splat. In the figure 5.3, we tried to render hairs with our technique we a simple hair as splat. We rotated it according to the normals and to increase the realism we changed the size of the splat depending on a procedural noise to have in some areas of the object smaller splat and in some others areas bigger splat. In our work, we did not focus on the color of rendered color because it depends on the 3d model and because our solution tends to be integrated into more complex pipeline rendering where the color is computed before.

Animations

During our analysis of the state of the art, we talked about what happened if the camera moves. We made some videos to show the results of our solution with an animated camera. The videos are in the Maxime's website (<http://www.maximeisnel.fr/procedural-stylization/>).

It is difficult to create a good *flatness* effect while moving, the animations of the brush painting is a good example at each frame we have something that has good *flatness* but during the motion we can see problem of *temporal continuity*. This due to the alpha blending and also mainly due to the size of the splat because they are big and so to improve *temporal continuity* we have to make to splat disappear/appear more slowly in order resolve some popping effect of the splats. As you can see in the animations of hairs rendering the splats are smaller so we don't have this problem of *temporal continuity*.

Performance

Our pipeline is implemented as a set of (non-optimized) OpenGL fragment shader in the Gratin software ([28]). Its performance is affected by the number of sampling during the creation of the procedural noises to minimize aliasing problem (popping effects). But the pipeline is mainly affected by the splatting step and the algorithm of order-independent transparency because in order to stylize the whole object we draw a lot of splat on the image so we have many splat cover others and then in order to know the order of them depending on the depth we sort with a *bubble sort* an array that contains at most 300 elements (arbitrary limits fixed depending on the GPU used) for each pixel of the final image. In other words, the performance depends on how many splats are drawn and their size, if the size of splat is big the splats will cover more splats than if their size is small and the complexity will be higher. For the animation of brush painting (present in Maxime's website) with 300 frames with a resolution of 512x512 pixels, it takes approximately 4 minutes with an Nvidia GeForce RTX 2070 graphics card to render.

— 6 —

Discussion and conclusion

6.0.1 Discussion

Watercolorization. We tried to render images imitating the effect of watercolor but we think that it needs to create procedural marks and to have more control on the transparency of these splats which is possible with some modifications in our method.

Depth in the splat itself. The splats are 2d images so when we draw them on the screen they keep their 2d aspect. But with the arrays needed for the *order-independent transparency* we can control the depth that we put in these arrays. So we imagine that we can make volumic splat (3d splats) in order to make for example hairs that intertwine.

Level of details. In our method, we use a global procedural noise to control the anchor points of the marks but we could combine more procedural noises in order to add more details in some parts of the object or as an artist will do add more details in the silhouettes.

6.0.2 Conclusion

We presented a post-processing stylization pipeline that draws marks over the screen while keeping a good *motion coherence*. Our method allows the user to choose the density of marks, the content of marks, their orientation and their size. Our main contribution in this report is the work on the Worley noise that permits to have a controllable texture which is getting closer to a point distribution with control of the size of points. The fractalization of the noise is not our contribution but we adapted it to our wanted, adding the choice of how many noises we mix and with which function interpolate the value. This solution only requires G-buffer data allowing it to be embedded in the compositing stage of most rendering software. This technique was then put to the test in order to make several images. We intend to reproduce existing styles which works with some, but some need modifications in our pipeline that can be done in future works.

Bibliography

- [1] *NPAR '00: Proceedings of the 1st International Symposium on Non-photorealistic Animation and Rendering*, New York, NY, USA, 2000. ACM. ACM Order No: 434000.
- [2] Loving vincent, 2017.
- [3] Alexandre Bléron, Romain Vergne, Thomas Hurtut, and Joëlle Thollot. Motion-coherent stylization with screen-space image filters. In *Proceedings of the Joint Symposium on Computational Aesthetics and Sketch-Based Interfaces and Modeling and Non-Photorealistic Animation and Rendering - Expressive '18*, pages 1–13, Victoria, British Columbia, Canada, 2018. ACM Press.
- [4] Adrien Bousseau, Fabrice Neyret, Joëlle Thollot, and David Salesin. Video watercolorization using bidirectional texture advection. *ACM Transactions on Graphics*, 26(3):104, July 2007.
- [5] Simon Breslav, Karol Szerszen, Lee Markosian, Pascal Barla, and Joëlle Thollot. Dynamic 2d Patterns for Shading 3d Scenes. page 6.
- [6] P. Bénard, A. Lagae, P. Vangorp, S. Lefebvre, G. Drettakis, and J. Thollot. A Dynamic Noise Primitive for Coherent Stylization. *Computer Graphics Forum*, 29(4):1497–1506, August 2010.
- [7] Pierre Bénard, Adrien Bousseau, and Joëlle Thollot. Dynamic solid textures for real-time coherent stylization. In *Proceedings of the 2009 symposium on Interactive 3D graphics and games - I3D '09*, page 121, Boston, Massachusetts, 2009. ACM Press.
- [8] Pierre Bénard, Adrien Bousseau, and Joëlle Thollot. State-of-the-Art Report on Temporal Coherence for Stylized Animations. *Computer Graphics Forum*, 30(8):2367–2386, December 2011.
- [9] Pierre Bénard, Jingwan Lu, Forrester Cole, Adam Finkelstein, and Joëlle Thollot. Active Strokes: Coherent Line Stylization for Animated 3d Models. page 10.
- [10] Ming-Te Chi and Tong-Yee Lee. Stylized and abstract painterly rendering system using a multiscale segmented sphere hierarchy. *IEEE transactions on visualization and computer graphics*, 12(1):61–72, February 2006.

- [11] Matthieu Cunzi, Joelle Thollot, Jean-Dominique Gascuel, Sylvain Paris, and Gilles Debunne. Dynamic Canvas for Non-Photorealistic Walkthroughs. page 10.
- [12] Eric Daniels. Deep canvas in disney's tarzan. In *ACM SIGGRAPH 99 Electronic Art and Animation Catalog*, SIGGRAPH '99, pages 124–, New York, NY, USA, 1999. ACM.
- [13] Bert Freudenberg, Maic Masuch, and Thomas Strothotte. Walk-Through Illustrations: Frame-Coherent Pen-and-Ink Style in a Game Engine. *Computer Graphics Forum*, 20(3):184–192, 2001.
- [14] Jennifer Fung and Oleg Veyrovka. "Pen-and-ink textures for real-time rendering". page 6.
- [15] Stéphane Grabli, Emmanuel Turquin, Frédo Durand, and François X. Sillion. Programmable rendering of line drawing from 3d scenes. *ACM Transactions on Graphics*, 29(2):1–20, March 2010.
- [16] Allison W. Klein, Wilmot Li, Michael M. Kazhdan, Wagner T. Corrêa, Adam Finkelstein, and Thomas A. Funkhouser. Non-photorealistic virtual environments. In *Proceedings of the 27th annual conference on Computer graphics and interactive techniques - SIGGRAPH '00*, pages 527–534, Not Known, 2000. ACM Press.
- [17] Yunjin Lee, Lee Markosian, Seungyong Lee, and John F Hughes. Line drawings via abstracted shading. page 5.
- [18] Liang Lin, Kun Zeng, Yizhou Wang, Ying-Qing Xu, and Song-Chun Zhu. Video Stylization: Painterly Rendering and Optimization with Content Extraction. *IEEE TRANSACTIONS ON CIRCUITS AND SYSTEMS FOR VIDEO TECHNOLOGY*, page 13.
- [19] Barbara J. Meier. Painterly rendering for animation. In *Proceedings of the 23rd annual conference on Computer graphics and interactive techniques - SIGGRAPH '96*, pages 477–484, Not Known, 1996. ACM Press.
- [20] Santiago E. Montesdeoca, Hock Soon Seah, Amir Semmo, Pierre Bénard, Romain Vergne, Joëlle Thollot, and Davide Benvenuti. MNPR: a framework for real-time expressive non-photorealistic rendering of 3d computer graphics. In *Proceedings of the Joint Symposium on Computational Aesthetics and Sketch-Based Interfaces and Modeling and Non-Photorealistic Animation and Rendering - Expressive '18*, pages 1–11, Victoria, British Columbia, Canada, 2018. ACM Press.
- [21] Cedrick Münstermann, Stefan Krumpen, Reinhard Klein, and Christoph Peters. Moment-based order-independent transparency. *Proc. ACM Comput. Graph. Interact. Tech.*, 1(1):7:1–7:20, July 2018.
- [22] Fabrice Neyret. iMAGIS-GRAVIR† / IMAG-INRIA. page 8.
- [23] Ken Perlin. Improving Noise. In *Proceedings of the 29th Annual Conference on Computer Graphics and Interactive Techniques*, SIGGRAPH '02, pages 681–682, New York, NY, USA, 2002. ACM. event-place: San Antonio, Texas.

- [24] Emil Praun, Hugues Hoppe, Matthew Webb, and Adam Finkelstein. Real-time hatching. In *Proceedings of the 28th annual conference on Computer graphics and interactive techniques - SIGGRAPH '01*, page 581, Not Known, 2001. ACM Press.
- [25] Johannes Schmid, Martin Sebastian Senn, Markus Gross, and Robert W. Sumner. Over-Coat: an implicit canvas for 3d painting. In *ACM SIGGRAPH 2011 papers on - SIGGRAPH '11*, page 1, Vancouver, British Columbia, Canada, 2011. ACM Press.
- [26] Michael Schwarz and Marc Stamminger. On predicting visual popping in dynamic scenes. In *Proceedings of the 6th Symposium on Applied Perception in Graphics and Visualization, APGV '09*, pages 93–100, New York, NY, USA, 2009. ACM.
- [27] David Vanderhaeghe and John Collomosse. Stroke Based Painterly Rendering. In Paul Rosin and John Collomosse, editors, *Image and Video-Based Artistic Stylisation*, volume 42, pages 3–21. Springer London, London, 2013.
- [28] Romain Vergne and Pascal Barla. Designing Gratin, A GPU-Tailored Node-Based System. *Journal of Computer Graphics Techniques*, 4(4):54–71, 2015.
- [29] Romain Vergne, David Vanderhaeghe, Jiazhou Chen, Pascal Barla, Xavier Granier, and Christophe Schlick. Implicit Brushes for Stylized Line-based Rendering. *Computer Graphics Forum*, 30(2):513–522, April 2011.
- [30] Steven Yantis and John Jonides. Abrupt visual onsets and selective attention: Evidence from visual search. *Journal of experimental psychology. Human perception and performance*, 10:601–21, 11 1984.
- [31] Kun Zeng, Mingtian Zhao, Caiming Xiong, and Song-Chun Zhu. From image parsing to painterly rendering. *ACM Transactions on Graphics*, 29(1):1–11, December 2009.